

# Introduction to **Information Retrieval**

Introducing ranked retrieval

# This unit and repetition

---

- **Today:** IR, VSM, mini intro to word embeddings
- **Repetition of last unit: NLP Basics**
  - Preprocessing: **tokenization** – what is it?
  - Preprocessing: **stopword removal** – what is it?
  - Preprocessing: **lemmatization / stemming** – what is it?
  - **NLTK** included corpora, functions to make: collocations, frequency distributions, ...

# Repetition – unit 1

---

- What is **IR**? What **evaluation measures** did we discuss?
- What is a **term-document matrix**? How to do Boolean queries on it?
- What is an **inverted index**?
  - Which **runtime** do basic (Boolean) **queries** have on an inverted index?
- What is a **Biword-index**? Why would we want to use it? Why not?

# Repetition – unit 1

---

- What is a **positional index**?
- **Which queries** can we do with the positional index?

# Ranked retrieval

---

- Thus far, our queries have all been Boolean.
  - Documents either match or don't.
- Good for expert users with precise understanding of their needs and the collection.
  - Also good for applications: Applications can easily consume 1000s of results.
- Not good for the majority of users.
  - Most users **incapable of writing Boolean queries** (or they are, but they think it's too much work).
  - Most users don't want to wade through 1000s of results.
    - This is particularly true of web search.

# Problem with Boolean search: feast or famine

---

- Boolean queries often result in either too few ( $\approx 0$ ) or too many (1000s) results.
  - Query 1: “*standard user dlink 650*”  $\rightarrow$  200,000 hits
  - Query 2: “*standard user dlink 650 no card found*”  $\rightarrow$  0 hits
- It takes a lot of skill to come up with a query that produces a manageable number of hits.
  - **AND** gives too few; **OR** gives too many

# Ranked retrieval models

---

- Rather than a set of documents satisfying a query expression, in **ranked retrieval models**, the system returns an **ordering over the (top) documents** in the collection with respect to a query
- **Free text queries:** Rather than a query language of operators and expressions, the user's query is just **one or more words in a human language**
- In principle, there are two separate choices here, but in practice, ranked retrieval models have normally been associated with free text queries and vice versa

# Feast or famine: not a problem in ranked retrieval

---

- When a system produces a ranked result set, large result sets are not an issue
  - Indeed, the **size of the result set is not an issue**
  - We just show the **top  $k$  ( $\approx 10$ )** results
  - We don't overwhelm the user
- Premise: the ranking algorithm works



# Scoring as the basis of ranked retrieval

---

- We wish to return in order the documents most likely to be useful to the searcher
- How can we rank-order the documents in the collection with respect to a query?
- **Assign a score – say in  $[0, 1]$  – to each document**
- This score measures how well document and query “match”.

# Query-document matching scores

---

- We need a way of assigning a score to a query/document pair
- Let's start with a **one-term** query
- If the query term **does not occur** in the document: score should be 0
- The **more frequent** the query term in the document, the **higher** the score (should be)
- We will look at a number of alternatives for this

# Introduction to **Information Retrieval**

Introducing ranked retrieval

# Introduction to **Information Retrieval**

Scoring with the Jaccard coefficient

# Take 1: Jaccard coefficient

---

- Simple **starting point for ranking**
- A **commonly used measure** of overlap of two sets  $A$  and  $B$  is the Jaccard coefficient
- $\text{jaccard}(A,B) = |A \cap B| / |A \cup B|$
- $\text{jaccard}(A,A) = 1$
- $\text{jaccard}(A,B) = 0$  if  $A \cap B = 0$
- $A$  and  $B$  don't have to be the same size.
- **Always assigns a number between 0 and 1.**

# Jaccard coefficient: Scoring example

---

- What is the query-document match score that the Jaccard coefficient computes for each of the two documents below?
- Query: *ides of march*
- Document 1: *caesar died in march*
- Document 2: *the long march*
- *Do  $Jaccard(q, d1)$  and  $Jaccard(q, d2)$  – what are the results*

# Exercise in Class (?)

---

- Implement Jaccard coefficient for some dataset which is short and simple (eg. first 100 words of documents in gutenberg NLTK corpus)
- Can also be used for **document similarity**

# Issues with Jaccard for scoring

---

- It doesn't consider *term frequency* (how many times a term occurs in a document)
  - **Rare terms** in a collection are more informative than frequent terms
  - Jaccard doesn't consider this information
- We need a more sophisticated way of **normalizing** for length
  - Later in this lecture, we'll use  $|A \cap B| / \sqrt{|A \cup B|}$  ... instead of  $|A \cap B| / |A \cup B|$  (Jaccard) for length normalization (as used in cosine similarity (later))



# Introduction to **Information Retrieval**

Scoring with the Jaccard coefficient

# Introduction to **Information Retrieval**

Term frequency weighting

# Recall: Binary term-document incidence matrix

---

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
Calpurnia	0	1	0	0	0	0
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0

Each document is represented by a *binary vector*  $\in \{0,1\}^{|V|}$

# Term-document count matrices

- Consider the **number of occurrences** of a term in a document:
  - Each document is a count vector in  $\mathbb{N}^{|V|}$ : a column below

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	157	73	0	0	0	0
Brutus	4	157	0	1	0	0
Caesar	232	227	0	2	1	1
Calpurnia	0	10	0	0	0	0
Cleopatra	57	0	0	0	0	0
mercy	2	0	3	5	5	1
worser	2	0	1	1	1	0

# Bag of words model

---

- Vector representation **doesn't consider the ordering** of words in a document
- *John is quicker than Mary* and *Mary is quicker than John* **have the same vectors**
- This is called the **bag of words** model.
- **In a sense, this is a step back:** The positional index was able to distinguish these two documents
  - We will look at “recovering” positional information later on
  - **For now: bag of words model**

# Term frequency tf

- The term frequency  $tf_{t,d}$  of term  $t$  in document  $d$  is defined as the number of times that  $t$  occurs in  $d$ .
- We want to use tf when computing query-document match scores. But how?
- **Raw term frequency** is not what we want:
  - A document with 10 occurrences of the term is more relevant than a document with 1 occurrence of the term.
  - But **not 10 times more relevant**.
- Relevance does **not increase proportionally** (linearly) with term frequency.

NB: frequency = count in IR

# Log-frequency weighting

- The **log frequency** weight of term  $t$  in  $d$  is

$$w_{t,d} = \begin{cases} 1 + \log_{10} \text{tf}_{t,d}, & \text{if } \text{tf}_{t,d} > 0 \\ 0, & \text{otherwise} \end{cases}$$

- $0 \rightarrow 0, 1 \rightarrow 1, 2 \rightarrow 1.3, 10 \rightarrow 2, 1000 \rightarrow 4$
- **Score for a document-query pair:** sum over terms  $t$  in both  $q$  and  $d$ , *only for intersecting terms*:
- **score**  $= \sum_{t \in q \cap d} (1 + \log \text{tf}_{t,d})$
- The score is **0** if **none** of the query terms is present in the document.

# Exercise TDM with log10-tf

---

- For gutenber corpus, create a TDM, with log-frequency weighting.  
(To make it faster, again use only the first 100 words per document).



# Introduction to **Information Retrieval**

Term frequency weighting

# Introduction to **Information Retrieval**

(Inverse) Document frequency weighting

# Document frequency

---

- **Rare terms** are more informative than frequent terms
  - Recall **stop words**
- Consider a term in the query that is rare in the collection (e.g., *arachnocentric*)
- A document containing this term is very likely to be relevant to the query *arachnocentric*
- → We want a **high weight for rare terms** like *arachnocentric*.

# Document frequency, continued

---

- **Frequent terms are less informative** than rare terms
- Consider a query term that is frequent in the collection (e.g., *high*, *increase*, *line*)
- A document containing such a term is more likely to be relevant than a document that doesn't
- But it's not a sure indicator of relevance.
- → For frequent terms, we want positive weights for words like *high*, *increase*, and *line*
- **But lower weights than for rare terms.**
- We will use **document frequency (df)** to capture this.

# idf weight

---

- $df_t$  is the document frequency of  $t$ : the number of documents that contain  $t$ 
  - $df_t$  is an **inverse measure** of the informativeness of  $t$
- We define the idf (inverse document frequency) of  $t$  by
$$idf_t = \log_{10} (N/df_t)$$
  - We use  $\log (N/df_t)$  instead of  $N/df_t$  to “**dampen**” the effect of idf.

Will turn out the base of the log is immaterial.

# idf example, suppose $N = 1\text{M}$

term	$df_t$	$idf_t$
calpurnia	1	$\rightarrow 6$
animal	100	?
sunday	1,000	?
fly	10,000	?
under	100,000	?
the	1,000,000	?

$$idf_t = \log_{10} (N/df_t)$$

There is **one idf value for each** term  $t$  in a collection.

# Effect of idf on ranking

---

- Question: Does **idf** have an effect on ranking for one-term queries, like
  - iPhone

# Effect of idf on ranking

---

- Question: Does idf have an effect on ranking for one-term queries, like
  - iPhone
- idf has **no effect on ranking one term queries**
  - idf affects the ranking of documents for queries with at **least two terms**
  - For the query **capricious person**, idf weighting makes occurrences of **capricious** count for much more in the final document ranking than occurrences of **person**.



# Collection vs. Document frequency

- The **collection frequency** of  $t$  is the number of occurrences of  $t$  in a collection, counting multiple occurrences. **Usually not used in IR, because:**
- Example:

Word	Collection frequency	Document frequency
<i>insurance</i>	10440	3997
<i>try</i>	10422	8760

- Which word is a **better search term** (and should get a higher weight)?

# Introduction to **Information Retrieval**

(Inverse) Document frequency weighting

# Introduction to **Information Retrieval**

tf-idf weighting

# tf-idf weighting

---

- The **tf-idf** weight of a term is the product of its tf weight and its idf weight.

$$w_{t,d} = (1 + \log \text{tf}_{t,d}) \times \log_{10}(N / \text{df}_t)$$

- **Best known weighting scheme in information retrieval**
  - Note: the “-” in tf-idf is a hyphen, not a minus sign!
  - Alternative names: tf.idf, tf x idf
- Increases with the **number of occurrences** within a document
- **Increases with the rarity of the term in the collection**

# Final ranking of documents for a query

---

$$\text{Score}(q, d) = \sum_{t \in q \cap d} \text{tf} \cdot \text{idf}_{t,d}$$

# Binary $\rightarrow$ count $\rightarrow$ weight matrix

---

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	5.25	3.18	0	0	0	0.35
Brutus	1.21	6.1	0	1	0	0
Caesar	8.59	2.54	0	1.51	0.25	0
Calpurnia	0	1.54	0	0	0	0
Cleopatra	2.85	0	0	0	0	0
mercy	1.51	0	1.9	0.12	5.25	0.88
worser	1.37	0	0.11	4.15	0.25	1.95

Each document is now represented by a real-valued vector of tf-idf weights  $\in \mathbb{R}^{|V|}$

# Exercise

---

- Extend TDM from previous exercise to have tf-idf weights
- You can base your code on `code9_tdm.py`

# Introduction to **Information Retrieval**

tf-idf weighting



# Repetition of unit 2

---

- What is / how?
  - Document ranking
  - Term frequency
  - Inverted document frequency
  - Tf-idf

# Term-Doc Matrix Creation

---

- Look at sample code:

`code10_tdm.py`

# This Unit: Overview

---

- Vector Space Model
- Term-term Matrix
- Word Embeddings

# Introduction to **Information Retrieval**

The Vector Space Model (VSM)

# Documents as vectors

---

- Now we have a  $|V|$ -dimensional vector space
- Terms are axes of the space
- Documents are points or vectors in this space
- Very high-dimensional: tens of millions of dimensions when you apply this to a web search engine
- These are **very sparse vectors** – most entries are zero

# Queries as vectors

---

- Key idea 1: Do the same for queries: represent them as vectors in the space
- Key idea 2: Rank documents according to their **proximity to the query** in this space
- proximity = similarity of vectors
- proximity  $\approx$  inverse of distance
- **Reason: We do this because we want to get away from the you're-either-in-or-out Boolean model**
- Instead: **rank** more relevant documents higher than less relevant documents


# Formalizing vector space proximity

---

- First cut: distance between two points
  - (= distance between the end points of the two vectors)
- **Euclidean distance?**
- Euclidean distance is a **bad idea** . . .
- . . . because Euclidean distance is **large** for vectors of **different lengths**.

# Why distance is a bad idea

The Euclidean distance between  $\vec{q}$  and  $\vec{d}_2$  is large even though the distribution of terms in the query  $\vec{q}$  and the distribution of terms in the document  $\vec{d}_2$  are very similar.





# Use angle instead of distance

---

- Thought experiment: take a document  $d$  and append it to itself. Call this document  $d'$ .
- “Semantically”  $d$  and  $d'$  have the same content
- The Euclidean distance between the two documents can be quite large
- The **angle between the two documents** is 0, corresponding to maximal similarity.
- Key idea: Rank documents according to angle with query.


# From angles to cosines

---

- The following two notions are equivalent.
  - Rank documents in decreasing order of the angle between query and document
  - Rank documents in increasing order of  $\text{cosine}(\text{query}, \text{document})$
- Cosine is a **monotonically decreasing function** for the interval  $[0^\circ, 180^\circ]$

# From angles to cosines

---



- But how – *and why* – should we be computing cosines?

# Length normalization

---

- A vector can be **(length-) normalized** by dividing each of its components by its length – for this we use the

**$L_2$  norm:**

$$\|\vec{x}\|_2 = \sqrt{\sum_i x_i^2}$$

- Dividing a vector by its  $L_2$  norm makes it a **unit (length) vector** (on surface of unit hypersphere)
- Effect on the two documents  $d$  and  $d'$  ( $d$  appended to itself) from earlier slide: they have identical vectors after length-normalization.
  - **Long and short documents now have comparable weights**

# cosine(query,document)

Dot product

Unit vectors

$$\cos(\vec{q}, \vec{d}) = \frac{\vec{q} \bullet \vec{d}}{|\vec{q}| |\vec{d}|} = \frac{\vec{q}}{|\vec{q}|} \bullet \frac{\vec{d}}{|\vec{d}|} = \frac{\sum_{i=1}^{|\mathcal{V}|} q_i d_i}{\sqrt{\sum_{i=1}^{|\mathcal{V}|} q_i^2} \sqrt{\sum_{i=1}^{|\mathcal{V}|} d_i^2}}$$

$q_i$  is the tf-idf weight of term  $i$  in the query

$d_i$  is the tf-idf weight of term  $i$  in the document

$\cos(\vec{q}, \vec{d})$  is the cosine similarity of  $\vec{q}$  and  $\vec{d}$  ... or,  
equivalently, the cosine of the angle between  $\vec{q}$  and  $\vec{d}$ .

# Cosine for length-normalized vectors


- For **length-normalized vectors**, cosine similarity is simply the **dot product** (or scalar product):

$$\cos(\vec{q}, \vec{d}) = \vec{q} \bullet \vec{d} = \sum_{i=1}^{|V|} q_i d_i$$

for  $q, d$  length-normalized.

(→ Length-normalize document vectors in advance!)

# Cosine similarity illustrated



# Cosine similarity amongst 3 documents

How similar are  
the novels

**SaS**: *Sense and  
Sensibility*

**PaP**: *Pride and  
Prejudice*, and

**WH**: *Wuthering  
Heights*?

term	SaS	PaP	WH
affection	115	58	20
jealous	10	7	11
gossip	2	0	6
wuthering	0	0	38

Term frequencies (counts)

Note: To simplify this example, we don't do idf weighting.



# Exercise VSM

---

- Given the frequency counts of the SaS, PaP, WH books:
  - Put them into vectors (lists)
  - Compute log-frequency
  - Create functions for L2-norm and dot-product
  - Apply L2-norm to vectors
  - Compute similarity between documents, print it.
  - Which 2 documents are most similar?

# 3 documents example contd.

## Log frequency weighting

term	SaS	PaP	WH
affection	3.06	2.76	2.30
jealous	2.00	1.85	2.04
gossip	1.30	0	1.78
wuthering	0	0	2.58

## After length normalization

term	SaS	PaP	WH
affection	0.789	0.832	0.524
jealous	0.515	0.555	0.465
gossip	0.335	0	0.405
wuthering	0	0	0.588

$$\cos(\text{SaS}, \text{PaP}) \approx$$

$$0.789 \times 0.832 + 0.515 \times 0.555 + 0.335 \times 0.0 + 0.0 \times 0.0 \approx 0.94$$

$$\cos(\text{SaS}, \text{WH}) \approx 0.79$$

$$\cos(\text{PaP}, \text{WH}) \approx 0.69$$

Why do we have  $\cos(\text{SaS}, \text{PaP}) > \cos(\text{SAS}, \text{WH})$ ?

# Introduction to **Information Retrieval**

Evaluating search engines

# Measures for a search engine

---

- How fast does it index
  - Number of documents/hour
  - (Average document size)
- How fast does it search
  - Latency as a function of index size
- Expressiveness of query language
  - Ability to express complex information needs
  - Speed on complex queries
- Uncluttered UI
- Is it free?

# Measures for a search engine

---

- All of the preceding criteria are *measurable*: we can quantify speed/size
  - we can make expressiveness precise
- The key measure: user happiness
  - What is this?
  - Speed of response/size of index are factors
  - But blindingly fast, useless answers won't make a user happy
- Need a way of quantifying user happiness with the results returned
  - Relevance of results to user's information need

# Evaluating an IR system

---

- An **information need** is translated into a **query**
- Relevance is assessed relative to the **information need** *not* the **query**
- E.g., Information need: *I'm looking for information on whether drinking red wine is more effective at reducing your risk of heart attacks than white wine.*
- Query: **wine red white heart attack effective**
- You evaluate whether the doc addresses the information need, not whether it has these words

# Evaluating ranked results

---

- Evaluation of a result set:
  - If we have
    - a benchmark document collection
    - a benchmark set of queries
    - assessor judgments of whether documents are relevant to queries

Then we can use Precision/Recall/F measure as before

- Evaluation of ranked results:
  - The system can return any number of results
  - By taking various numbers of the top returned documents (levels of recall), the evaluator can produce a *precision-recall curve*

# Recall/Precision

---

R      P

- 1 R
- 2 N
- 3 N
- 4 R
- 5 R
- 6 N
- 7 R
- 8 N
- 9 N
- 10 N

Assume 10 rel docs  
in collection



# Two current evaluation measures...

---

- Mean average precision (MAP)
  - AP: Average of the precision value obtained for the top  $k$  documents, each time a relevant doc is retrieved
  - Avoids interpolation, use of fixed recall levels
  - Does weight most accuracy of top returned results
  - MAP for set of queries is arithmetic average of APs
    - Macro-averaging: each query counts equally

# Introduction to **Information Retrieval**

Evaluating search engines

# Term-Term Matrices

---

- In IR systems we typically use Term-Document matrices
- But in many NLP applications we are interested in **term-term co-occurrence** matrices
- Co-occurrence can be measured in different ways, for example
  - Within a unit like a **sentence** or **paragraph**
  - Within a **word window** (left and/or right) of the target word – eg. a word window of 5
  - **Q: when could a co-occurrence matrix be useful?**

# Co-occurrence matrix - Example

1. I enjoy flying.
2. I like NLP.
3. I like deep learning.

The resulting counts matrix will then be:

$$X = \begin{array}{c} \begin{array}{c} I \\ like \\ enjoy \\ deep \\ learning \\ NLP \\ flying \\ . \end{array} \begin{bmatrix} I & like & enjoy & deep & learning & NLP & flying & . \\ 0 & 2 & 1 & 0 & 0 & 0 & 0 & 0 \\ 2 & 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 1 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 & 1 \\ 0 & 1 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 1 & 1 & 1 & 0 \end{bmatrix} \end{array}$$

# Co-occurrence Matrices

---

- Distributional hypothesis – "a word is characterized by the company it keeps"
- Words are defined by their context (words)

A bottle of **tesgüino** is on the table  
Everybody likes **tesgüino**  
**Tesgüino** makes you drunk  
We make **tesgüino** out of corn.

# Co-occurrence Matrices

---

- These matrices are often input to further processing, eg. for word embeddings, SVD, etc.
- Weighting of the matrix:
  - Raw counts: but not best option
    - “the” and “of” are very frequent, but maybe not the most discriminative
  - PPMI (see next slide)
    - Goal: context is informative about the target word
  - ...

# Co-occurrence Matrices – PPMI

---

- Starts from co-occurrence counts

## PMI between two words: (Church & Hanks 1989)

Do words  $x$  and  $y$  co-occur more than if they were independent?

$$\text{PMI}(\text{word}_1, \text{word}_2) = \log_2 \frac{P(\text{word}_1, \text{word}_2)}{P(\text{word}_1)P(\text{word}_2)}$$

- PMI goes from neg. inf. to pos. inf.
- **PPMI**: we just set all negative values to 0.
  - Low counts are unreliable, and neg. relation hard to understand for humans

# Co-occurrence Matrices

---

- Co-occ. matrices are sparse representations
- Next time we look at dense representations, like word embeddings



# Mini Intro to Language Models and Word Embeddings

---

- What are **language models**?
- What are **word embeddings**?
  - Dense, floating point
  - Popular esp since word2vec (Mikolov et al, 2013)
  - Based on distributional hypothesis (Harris, 1957)
  - What does a vector look like, what are characteristics?
  - Trained on large corpora
  - Different training algorithms (GloVe, word2vec, ..)
  - Applications: many, eg. clustering, input to ML/DL, almost NLP task as representations of words